

P A S S

PROFESSIONAL ASSOCIATION
FOR SQL SERVER

PASS Deutschland e.V.

**100% Verfügbarkeit bei 24 Stunden Wartungsfenster
pro Tag**

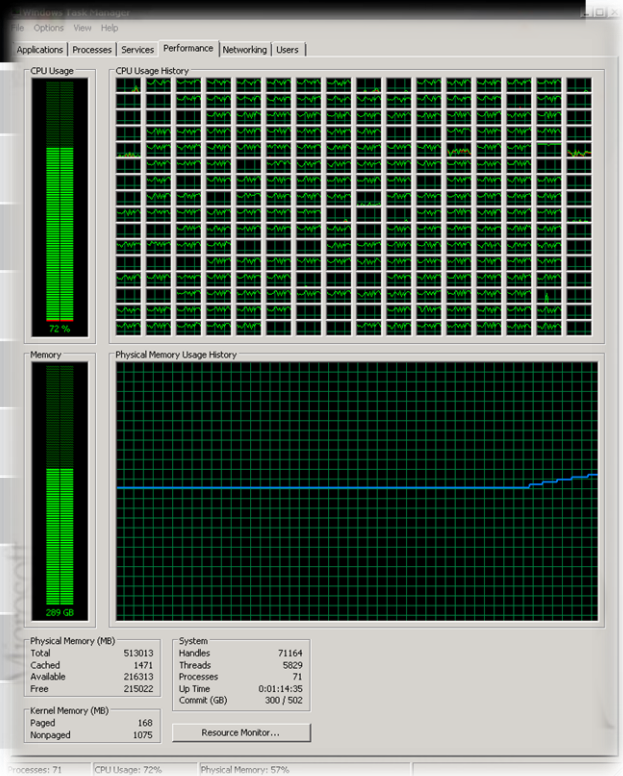
Original Title

100% Availability without 100% Uptime

=tg= Thomas Grohser

select * from =tg=

@@Version	Remark
SQL 4.21	First SQL Server ever used (1994)
SQL 6.0	First Log Shipping with failover
SQL 6.5	First SQL Server Cluster (NT4.0 + Wolfpack)
SQL 7.0	2+ billion rows / month in a single Table
SQL 2000	938 days with 100% availability
SQL 2000 IA64	First SQL Server on Itanium IA64
SQL 2005 IA64	First OLTP long distance database mirroring
SQL 2008 IA64	First Replication into mirrored databases
SQL 2008R2 IA64	First 256 CPUs & >500.000 STMT/sec
SQL 11 (Denali)	Can't wait to push the limits even further

A screenshot of the Windows Task Manager Performance tab. The CPU Usage section shows a green bar at 72% and a 'CPU Usage History' graph with a green grid. The Memory section shows a green bar at 269 GB and a 'Physical Memory Usage History' graph with a blue line. At the bottom, system statistics are listed: Physical Memory (MB) with Total 513013, Cached 1471, Available 216313, and Free 215022; Kernel Memory (MB) with Paged 168 and Nonpaged 1075; System Handles 71164, Threads 5829, Processes 71, Up Time 0:01:14:35, and Commit (GB) 300 / 502. A 'Resource Monitor...' button is also visible.

Thomas Grohser, SQL Server MVP, bwin Interactive Entertainment AG

<http://www.grohser.com/>

Focus on SQL Server Infrastructure Architecture and Implementation

Close Relationship with Microsoft

SQLCAT (SQL Server Customer Advisory Team)

SCAN (SQL Server Customer Advisory Network)

TAP (Technology Adoption Program SQL2008R2 and SQL11)

Active PASS member and PASS Summit Speaker





World's biggest publicly listed online gaming platform

World's leading provider of online Sports Betting

One of the largest **Poker networks**

Comprehensive range of **Payment Service Providing**

Integrated gaming portal - **22 languages,**
25 core markets

Gross gaming revenues 2008 (GGR):
EUR 421 million

More than **20 million registered customers**

1,500 employees

bwin builds on the strengths of the web in order
to **tie up responsibility and gaming**

15 million page views and up to **980,000**
users a day



Agenda

- Introduction to the challenge
- Theoretical solution
- Practical implementation
- Other problems we solved too
- Q+A



ATTENTION:
Important
Information may be
displayed at any
slide at any time!

The Challenge

- **Globally operating companies**
 - Global websites
 - Offices around the globe
- **They all face the challenge that there is no time left to perform service and maintenance on the database systems.**
 - More and more patches to apply
 - Less and less time to maintain


The Challenge

- **Every day of the year we must operate 24h per day without interruptions and do so faster and faster**
- **Every day of the year we would like to maintain our systems 24h per day.**

Non working ideas

- **Scale up**
 - Solves partially performance
- **Scale out by hash(UserID or UserName)**
 - Solves only availability and performance

**We need a glass that
can hold twice as much
content than it has volume...**



Solution

- Impossible to do with a technical solution?
- So we need a political solution

Yes we can

© Copyright and
™ Trademark by
some politician
from Hawaii ...
mahalo

Solution

- **What does a good politician do**
 - Bend the truth
 - and**
 - Let the people solve the problem.
- **And this is exactly what we will do...**

Solution

- **We promised a solution that is available 24 hours per day**
- **What does available mean?**
 - A system is considered available if **all users** of the system at any time they desire to use the system can do so successfully
 - We split the rule
 - A system is considered available by **a user** when the user of the system at any time he or she desires to use the system can do so successfully
 - If **each** individual **user** considers the system available the whole system is considered available

Solution

- **Facts**


- Our users are human beings
(People on the web or in an office)
- And all people have some basic needs
 - Eating
 - Sleeping
- The first one is not helping us much but the second ...
- Humans need at least 4 hours of sleep and most sleep more like 6 to 10 hours a day.
- And the greatest fact is that even if they travel the globe they somehow keep some of the hours of sleep and inactivity constant.

Solution

- So lets bend the facts into the rules
- If we group our users in a way that all people with the same inactivity pattern are in one group
- And place each group on its own server
- We just solved the problem...
- Easy isn't it?


Theoretical Solution

- Activity per hour of a single user on a single day




Theoretical Solution

- Activity per hour of a single user in one year



Theoretical Solution

- Activity per hour of multiple users in one year



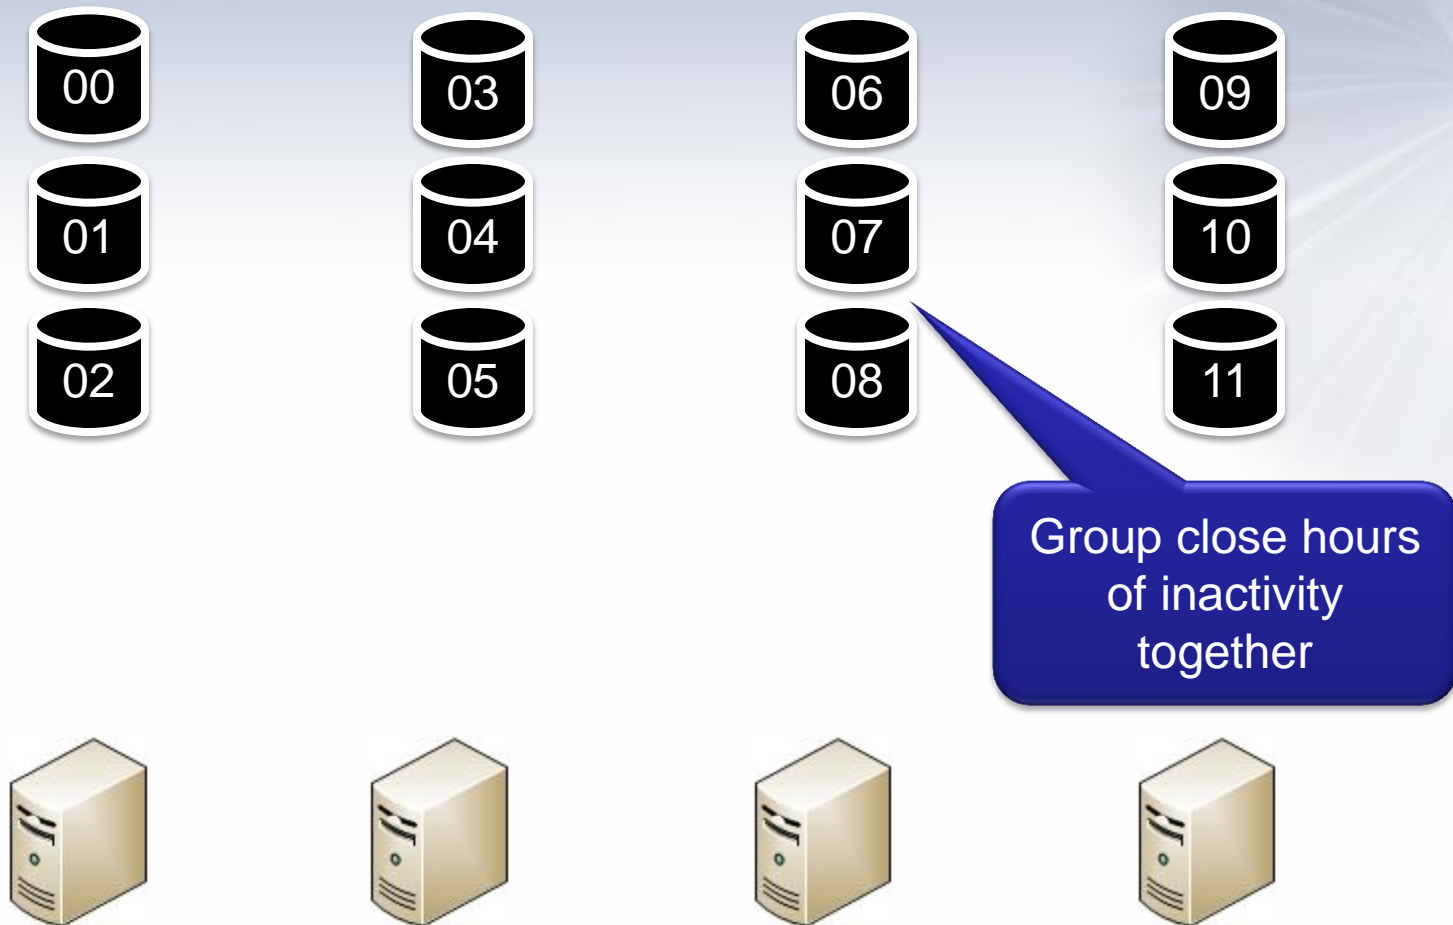
Some Statistics

- Large globally operating website with millions of users and several years of operational data to analyze
- 99.99% of the users have at least one constant inactive hour over all the years
- 99.99% of the users are have more than one constant inactive hours over the last 2 years
- 100% of the users have at least one constant inactive hour over the last year


Practical Solution

- **24 Partitions**
 - One for each hour of the day
 - 2 brick groups with 12 Partitions each
 - 4 bricks with 3 Partitions each


Why 12 Partitions on 4 Servers



Why 12 Partitions on 4 Servers




Why 12 Partitions on 4 Servers




Clustering or Mirroring/Log Shipping

- **Clustering**
 - Data only available once = Single point of failure
- **Mirroring/Log Shipping**
 - Out of the box not as automatic as clustering
 - Limited support for Replication

Cluster



Mirror & Log Shipping



Principal database


Mirrored database



Log shipped database

Log shipped database 2nd copy

Mirror & Log Shipping



Principal database


Mirrored database



Log shipped database

Log shipped database 2nd copy

Mirror & Log Shipping



Principal database


Mirrored database



Log shipped database

Log shipped database 2nd copy

Mirror & Log Shipping



Principal database

Mirrored database



Log shipped database

Log shipped database 2nd copy

Mirroring - Dual datacenter

- If you want the synchronous mirror to be always in the other datacenter it does not work with 4 servers and 12 partitions but works great with 8 Servers and 24 partitions
- Have fun figuring out the distribution matrix

How does it work



login


Connect to
db(4)

Work as
usual


Partition
ID=4

Login
Service


SRV1




SRV2



SRV3



SRV4



the Definitive Community for SQL Server Professionals

Join PASS Today at www.sqlpass.de

Other user



login

Connect to
db(1)


Work as
usual

Partition
ID=1


Login
Service




SRV1




SRV2



SRV3



SRV4




User sleeps...



db(4) no
activity

Rollout
Manager:
ALTER ...



Partition
ID=4

Login
Service

DBA:
REBUILD
INDEX

Join PASS The Definitive Community for SQL Server Professionals


Join PASS Today at www.sqlpass.de

Many users sleep...



db(3 & 4 & 5) no activity

DBA:
PATCH Server



Partition
ID=2,3,4,5

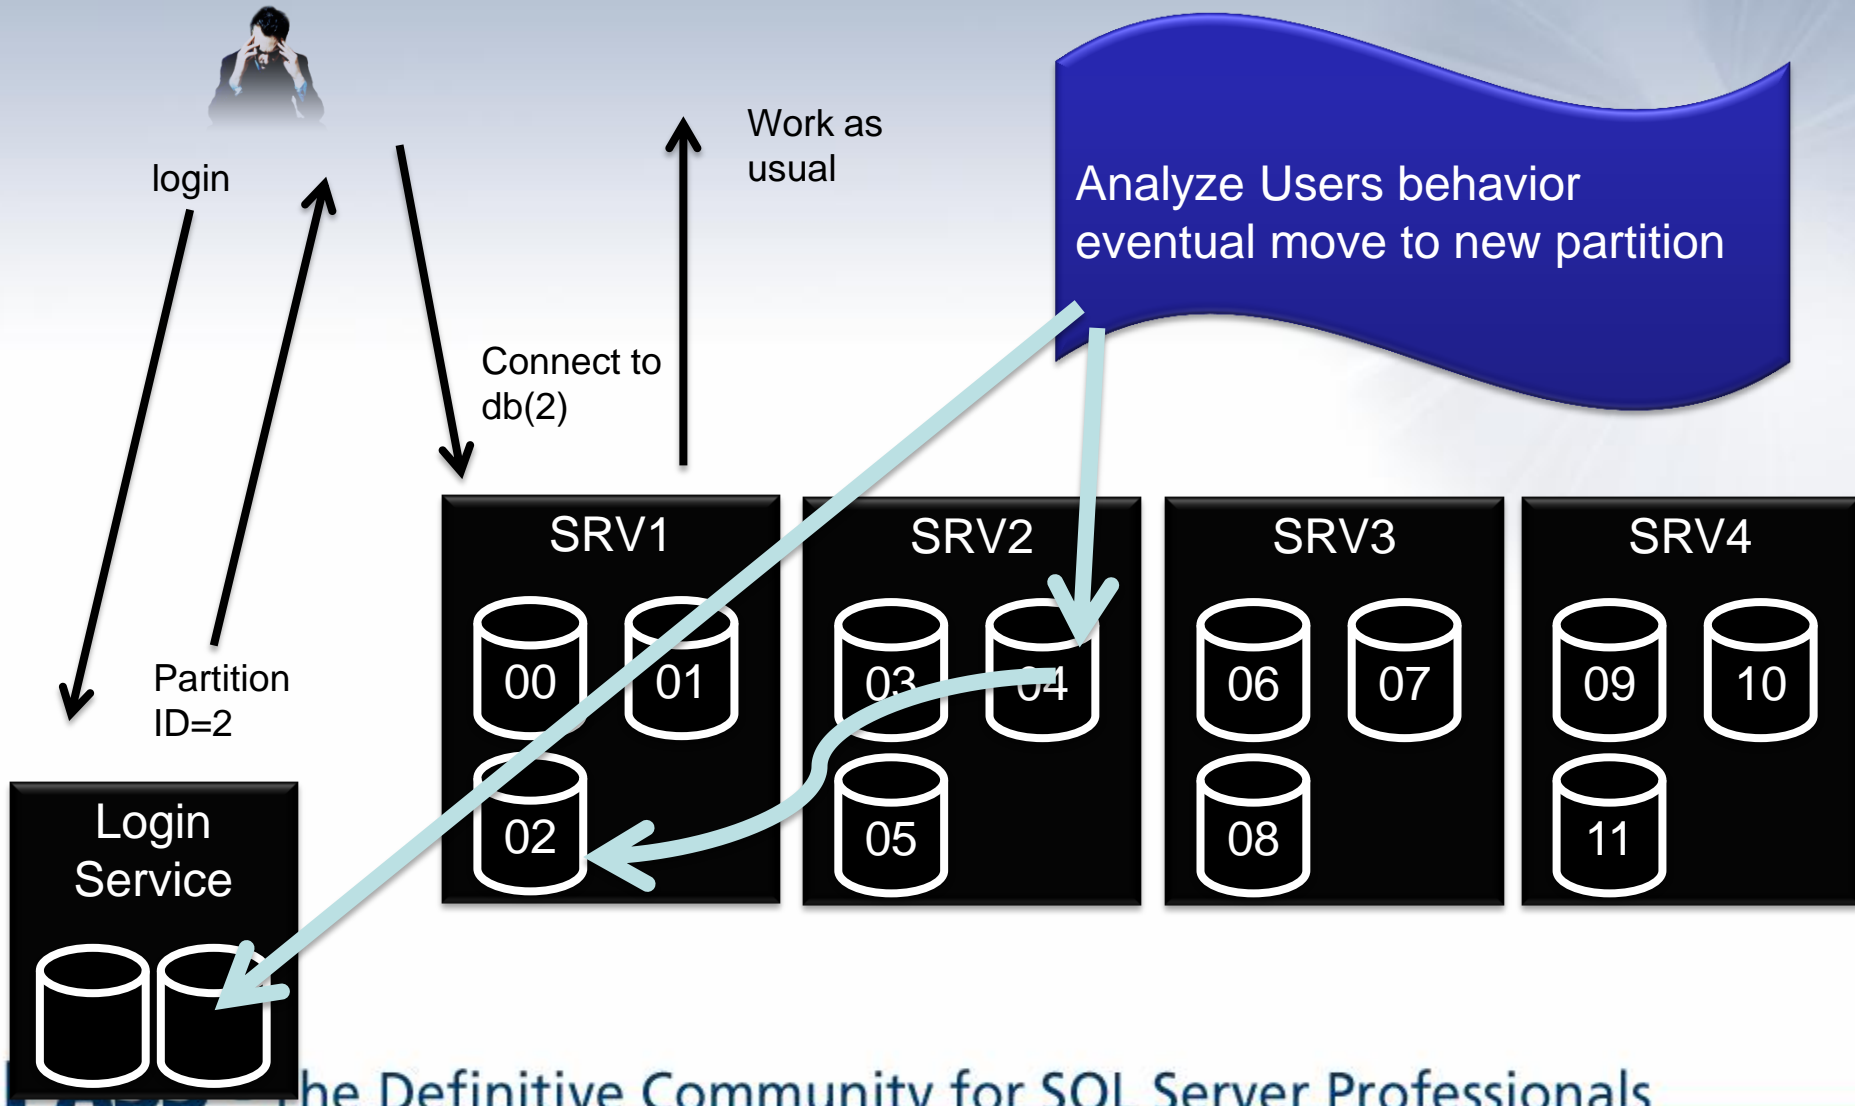
Login
Service



The Definitive Community for SQL Server Professionals

Join PASS Today at www.sqlpass.de

Repartition Users



Implementation Details

- **Identity Columns**

- GUID

- Large, very bad for indexes

- Ranges per partition

- 100.000.000 to 199.999.999 partition 1
 - 200.000.000 to 299.999.999 partition 2
 -
 - Large numbers from the beginning

- Renumber every time you move a customer

- Your DWH people will hate you...

Implementation Details

- **IDENTITY Columns**

- Reverse Ranges

- Partition 1

```
CREATE TABLE MyData  
(ID bigint IDENTITY(1, 65536))
```

- Partition 2

```
CREATE TABLE MyData  
(ID bigint IDENTITY(2, 65536))
```

- ...

Implementation Details

- **Global static data (Countries, ZipCodes, ...)**
 - Replicate/copy from one "master" source to each database or keep on separate server
- **Global dynamic data (ProductCatalog, ...)**
 - Keep on separate server or replicate/copy from one "master" source to each database
- **Customer data (Transactions, Invoices, ...)**
 - Partition with customer
- **Relationship data (FriendsOfFriends)**
 - Store on both ends


Other problems we solved too

- **Archiving**
- **Geographical placement of data**


Archiving

- Archiving OLTP data can be tricky and painful but management likes it because it saves money
- Create an extra Partition (the archive)
- Place the partition on a small server with large and inexpensive storage
- Move all inactive customers to this partition
- If the customer becomes active again move him back

Archiving



- High performance
- High availability
- “Low” Capacity

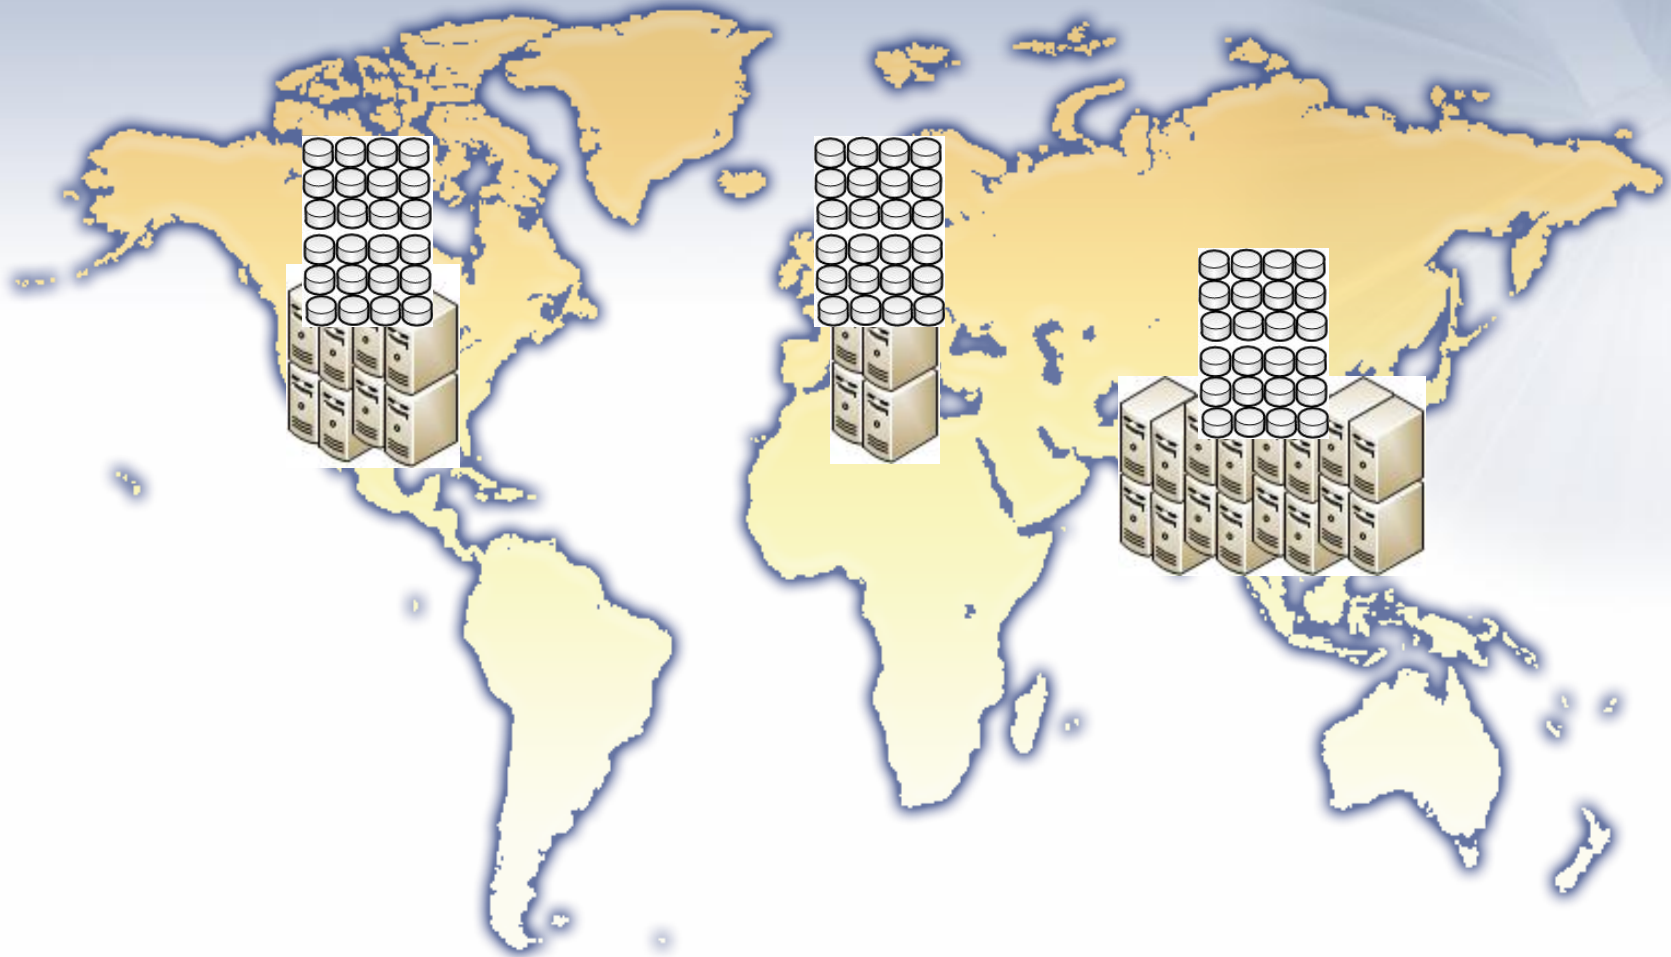


- Low performance
- Low availability
- High Capacity

Geographical placement

- Sometimes the requirement to place data at a specific location comes up.
- You can easily create extra partition(s) and place them on servers at a different location.
- For example: 24 Partitions on 4 Servers in Europe, 24 Partitions on 8 Servers in the US and 24 Partitions on 16 Servers in Asia

Geographical placement



PASS - The Definitive Community for SQL Server Professionals

Join PASS Today at www.sqlpass.de

Just some Ideas

- **Have a friendly user partition for new releases and features**
- **Let the user choose its hour of inactivity in a configuration page/dialog**

Wrap Up

- **Find a way to group your data in active/inactive portions over time**
- **Separate the data physically by this grouping**
- **Scale Out**
- **Gain manageability**
- **Gain availability**

Questions?

=tg= Thomas Grohser

tg@grohser.com

<http://www.grohser.com/>